

Human Memory Layer (HML) - Technical Documentation

Version 1.0.0 | January 2025

Table of Contents

- [1. Executive Summary](#)
 - [2. Architecture Overview](#)
 - [3. Core Protocols](#)
 - [4. Implementation Specifications](#)
 - [5. API Reference](#)
 - [6. Security & Privacy](#)
 - [7. Federation Network](#)
 - [8. Development Guide](#)
 - [9. Deployment Strategies](#)
 - [10. Appendices](#)
-

Executive Summary

The Human Memory Layer (HML) is a foundational infrastructure protocol that provides persistent, searchable, and transferable memory capabilities for AI agents and human users. It enables AI systems to maintain context across conversations, platforms, and time, while ensuring users retain complete ownership and control of their cognitive data.

Key Principles

- User Sovereignty:** Users own their memories completely
- Protocol, Not Platform:** Open standards enable interoperability
- Privacy by Design:** Local-first architecture with optional federation
- AI-Native:** Designed for seamless AI agent integration
- Decentralized:** No single point of failure or control

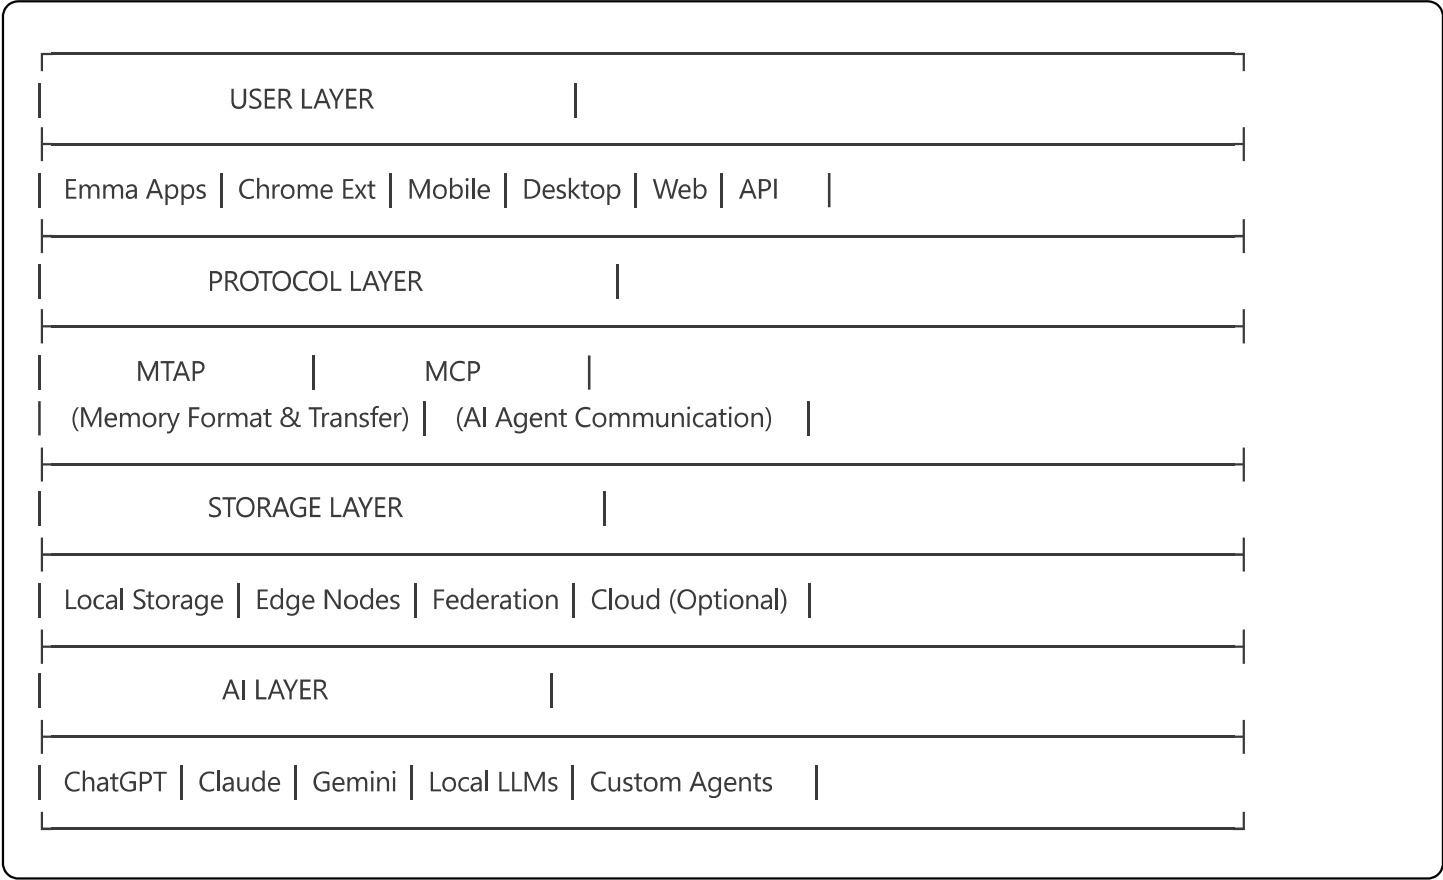
Core Components

- MTAP (Memory Transfer & Access Protocol):** Standardized memory format
- MCP (Model Context Protocol):** AI agent communication bridge

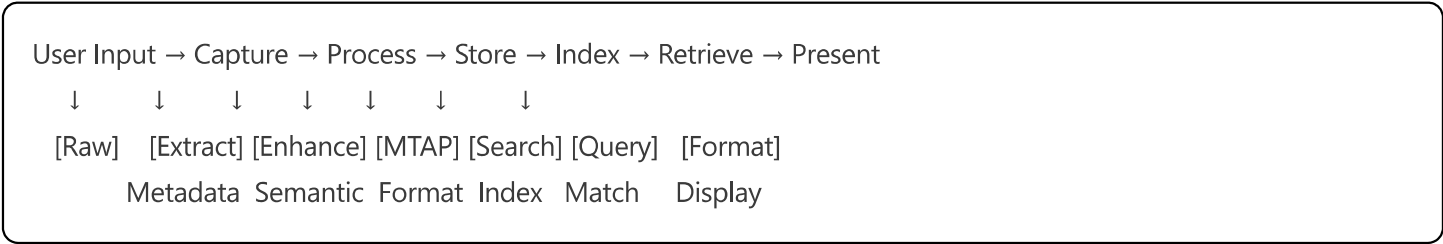
- **Emma**: Reference implementation and user-facing application
- **Federation Network**: Distributed memory sharing and replication

Architecture Overview

System Architecture



Data Flow Architecture



Memory Lifecycle

1. CREATION
- Input → Validation → Enrichment → Encoding → Storage
2. RETRIEVAL
- Query → Search → Rank → Filter → Format → Deliver
3. FEDERATION
- Local → Sync → Replicate → Verify → Distribute
4. EVOLUTION
- Access → Learn → Update → Version → Archive

Core Protocols

MTAP (Memory Transfer & Access Protocol)

Protocol Specification v1.0

json

```
{
  "mtap_version": "1.0.0",
  "memory": {
    "header": {
      "id": "uuid-v4",
      "version": "1.0.0",
      "created": "ISO-8601",
      "modified": "ISO-8601",
      "creator": "did:method:identifier",
      "signature": "base64-signature",
      "contentHash": "sha256-hash",
      "protocol": "MTAP/1.0"
    },
    "core": {
      "type": "text|image|audio|video|composite",
      "content": "actual-content-or-reference",
      "encoding": "UTF-8|base64|custom",
      "encrypted": false,
      "compression": null,
      "size": 1024
    },
    "semantic": {
      "summary": "AI-generated-summary",
      "keywords": ["keyword1", "keyword2"],
      "entities": [
        {
          "type": "person|place|org|date|url",
          "value": "entity-value",
          "confidence": 0.95
        }
      ],
      "emotions": ["joy", "curiosity"],
      "topics": ["AI", "memory"],
      "embeddings": {
        "model": "embedding-model-id",
        "vector": [0.1, 0.2, ...],
        "dimensions": 1536
      }
    },
    "relations": {
      "previous": "memory-id",
      "next": "memory-id",
      "parent": "memory-id",

```

```
"children": ["memory-id-1", "memory-id-2"],
"related": ["memory-id-3", "memory-id-4"],
"references": ["url1", "url2"],
"threads": ["thread-id-1"]
},
"permissions": {
  "owner": "did:method:identifier",
  "public": false,
  "shared": [
    {
      "did": "did:method:identifier",
      "permissions": ["read", "reference"],
      "expiry": "ISO-8601"
    }
  ],
  "agents": [
    {
      "agentId": "agent-identifier",
      "permissions": ["read", "write", "share"],
      "granted": "ISO-8601",
      "expiry": null
    }
  ],
  "encryption": {
    "method": "AES-256-GCM",
    "keyId": "key-identifier"
  }
},
"metadata": {
  "source": "chatgpt|claude|manual|api",
  "application": "emma-lite",
  "deviceId": "device-uuid",
  "location": {
    "lat": 0.0,
    "lon": 0.0,
    "accuracy": 10,
    "name": "location-name"
  },
  "context": {
    "activity": "conversation|research|note",
    "project": "project-id",
    "tags": ["tag1", "tag2"]
  },
  "custom": {}
}
```

```
},
"verification": {
  "witnesses": [
    {
      "did": "did:method:identifier",
      "signature": "base64-signature",
      "timestamp": "ISO-8601"
    }
  ],
  "blockchain": {
    "network": "ethereum|bitcoin|custom",
    "transaction": "tx-hash",
    "block": 12345678
  }
}
}
```

MTAP Operations

1. Create Memory

javascript

POST /mtap/memory

```
{
  "content": "Memory content",
  "type": "text",
  "metadata": {...}
}
```

Response:

```
{
  "id": "mem_abc123",
  "address": "mtap://memory/sha256hash",
  "created": "2025-01-20T10:00:00Z"
}
```

2. Retrieve Memory

javascript

```
GET /mtap/memory/{id}
GET mtap://memory/{contentHash}

Response: Full MTAP memory object
```

3. Update Memory

```
javascript

PATCH /mtap/memory/{id}
{
  "operations": [
    {"op": "add", "path": "/metadata/tags", "value": ["new-tag"]},
    {"op": "replace", "path": "/core/content", "value": "updated"}
  ]
}
```

4. Search Memories

```
javascript

POST /mtap/search
{
  "query": "search terms",
  "filters": {
    "type": "text",
    "dateRange": {
      "start": "2024-01-01",
      "end": "2024-12-31"
    },
    "source": ["chatgpt", "claude"]
  },
  "limit": 20,
  "offset": 0
}
```

MCP (Model Context Protocol)

Protocol Bridge Specification

```
javascript
```

// MCP Request Format

```
{
  "protocol": "MCP/1.0",
  "operation": "getContext",
  "parameters": {
    "query": "User intent or question",
    "maxTokens": 4000,
    "relevanceThreshold": 0.7,
    "timeRange": {
      "start": "ISO-8601",
      "end": "ISO-8601"
    },
    "sources": ["chatgpt", "claude", "manual"],
    "format": "markdown|json|xml"
  },
  "authentication": {
    "method": "bearer",
    "token": "jwt-token"
  }
}
```

// MCP Response Format

```
{
  "protocol": "MCP/1.0",
  "context": {
    "memories": [
      {
        "id": "memory-id",
        "content": "Memory content",
        "relevance": 0.95,
        "timestamp": "ISO-8601",
        "source": "chatgpt",
        "metadata": {...}
      }
    ],
    "summary": "Contextual summary",
    "tokens": 3500,
    "truncated": false
  },
  "performance": {
    "searchTime": 45,
    "processingTime": 120,
    "memoryCount": 15
  }
}
```



```
}  
}
```

MCP Integration Points

1. Direct API Access

```
python  
  
# Python SDK Example  
from hml import MCP  
  
mcp = MCP(api_key="your-api-key")  
context = mcp.get_context(  
    query="What did we discuss about quantum computing?",  
    max_tokens=2000  
)
```

2. LangChain Integration

```
python  
  
from langchain.memory import HMLMemory  
  
memory = HMLMemory(  
    api_key="your-api-key",  
    user_id="user-did"  
)  
  
chain = ConversationChain(  
    llm=llm,  
    memory=memory  
)
```

3. OpenAI Function Calling

```
javascript
```

```
{
  "name": "query_hml_memory",
  "description": "Search user's memory layer",
  "parameters": {
    "type": "object",
    "properties": {
      "query": {
        "type": "string",
        "description": "Search query"
      }
    }
  }
}
```

Implementation Specifications

Storage Implementations

1. Local Storage (Browser)

```
javascript
```

```
// IndexedDB Schema
```

```
{  
  databases: {  
    EmmaLiteDB: {  
      version: 2,  
      stores: {  
        memories: {  
          keyPath: "id",  
          indexes: ["timestamp", "source", "type"]  
        },  
        mtap_memories: {  
          keyPath: "header.id",  
          indexes: ["header.created", "header.contentHash"]  
        },  
        embeddings: {  
          keyPath: "memoryId",  
          indexes: ["timestamp"]  
        }  
      }  
    }  
  }  
}
```

2. Edge Storage (Cloudflare Workers)

```
javascript
```

```
// Durable Objects Schema
```

```
class MemoryStore {  
  constructor(state, env) {  
    this.state = state;  
    this.storage = state.storage;  
  }  
  
  async store(memory) {  
    const key = `mem:${memory.header.id}`;  
    await this.storage.put(key, memory);  
    await this.updateIndex(memory);  
  }  
  
  async retrieve(id) {  
    return await this.storage.get(`mem:${id}`);  
  }  
  
  async search(query, options) {  
    // Vector search implementation  
    const embeddings = await this.storage.list({  
      prefix: "emb:",  
      limit: 1000  
    });  
  
    return this.rankByRelevance(embeddings, query);  
  }  
}
```

3. Federation Storage (IPFS)

```
javascript
```

```
// IPFS Integration
import { create } from 'ipfs-http-client';

class IPFSMemoryStore {
  constructor() {
    this.ipfs = create({
      host: 'ipfs.infura.io',
      port: 5001,
      protocol: 'https'
    });
  }

  async store(memory) {
    const { cid } = await this.ipfs.add(
      JSON.stringify(memory)
    );

    await this.ipfs.pin.add(cid);
    return `ipfs://${cid}`;
  }

  async retrieve(cid) {
    const stream = this.ipfs.cat(cid);
    let data = '';

    for await (const chunk of stream) {
      data += chunk.toString();
    }

    return JSON.parse(data);
  }
}
```

Search & Retrieval

Vector Search Implementation

```
python
```

```
import numpy as np
from sentence_transformers import SentenceTransformer

class VectorSearch:
    def __init__(self):
        self.model = SentenceTransformer('all-MiniLM-L6-v2')
        self.index = None

    def create_embeddings(self, texts):
        return self.model.encode(texts)

    def build_index(self, memories):
        texts = [m['content'] for m in memories]
        embeddings = self.create_embeddings(texts)

        # Using FAISS for efficient similarity search
        import faiss
        dimension = embeddings.shape[1]
        self.index = faiss.IndexFlatL2(dimension)
        self.index.add(embeddings)

    def search(self, query, k=10):
        query_embedding = self.create_embeddings([query])
        distances, indices = self.index.search(query_embedding, k)
        return indices[0], distances[0]
```

Hybrid Search (Text + Vector)

```
javascript
```

```

class HybridSearch {
  async search(query, options = {}) {
    // Parallel search strategies
    const [textResults, vectorResults, graphResults] = await Promise.all([
      this.textSearch(query, options),
      this.vectorSearch(query, options),
      this.graphSearch(query, options)
    ]);

    // Fusion ranking
    return this.fuseResults([
      { results: textResults, weight: 0.3 },
      { results: vectorResults, weight: 0.5 },
      { results: graphResults, weight: 0.2 }
    ], options.limit);
  }

  fuseResults(resultSets, limit) {
    const scores = new Map();

    for (const { results, weight } of resultSets) {
      results.forEach((result, index) => {
        const score = (1 / (index + 1)) * weight;
        const existing = scores.get(result.id) || 0;
        scores.set(result.id, existing + score);
      });
    }

    return Array.from(scores.entries())
      .sort((a, b) => b[1] - a[1])
      .slice(0, limit)
      .map(([id, score]) => ({ id, score }));
  }
}

```

API Reference

REST API Endpoints

Authentication

http

POST /api/v1/auth/token

Content-Type: application/json

```
{
  "did": "did:web:user.example.com",
  "signature": "base64-signature",
  "timestamp": "ISO-8601"
}
```

Response: 200 OK

```
{
  "token": "jwt-token",
  "expires": "ISO-8601"
}
```

Memory Operations

Create Memory

http

POST /api/v1/memories

Authorization: Bearer {token}

Content-Type: application/json

```
{
  "content": "Memory content",
  "type": "text",
  "source": "api",
  "metadata": {
    "tags": ["important"],
    "project": "project-id"
  }
}
```

Response: 201 Created

```
{
  "id": "mem_abc123",
  "address": "mtap://memory/hash",
  "created": "ISO-8601"
}
```

Get Memory

http

GET /api/v1/memories/{id}

Authorization: Bearer {token}

Response: 200 OK

```
{  
  Full MTAP memory object  
}
```

Search Memories

http

POST /api/v1/memories/search

Authorization: Bearer {token}

Content-Type: application/json

```
{
  "query": "search terms",
  "filters": {
    "type": ["text", "image"],
    "source": ["chatgpt"],
    "dateRange": {
      "start": "2024-01-01",
      "end": "2024-12-31"
    }
  },
  "sort": {
    "field": "relevance|date",
    "order": "desc"
  },
  "pagination": {
    "limit": 20,
    "offset": 0
  }
}
```

Response: 200 OK

```
{
  "results": [...],
  "total": 145,
  "hasMore": true
}
```

WebSocket API

Real-time Memory Stream

javascript

```
const ws = new WebSocket('wss://api.emma-hml.com/v1/stream');

ws.on('open', () => {
  ws.send(JSON.stringify({
    type: 'authenticate',
    token: 'jwt-token'
  }));

  ws.send(JSON.stringify({
    type: 'subscribe',
    channels: ['memories', 'updates']
  }));
});

ws.on('message', (data) => {
  const event = JSON.parse(data);
  switch(event.type) {
    case 'memory.created':
      handleNewMemory(event.data);
      break;
    case 'memory.updated':
      handleUpdatedMemory(event.data);
      break;
  }
});
```

GraphQL API

graphql

Schema Definition

```
type Memory {  
  id: ID!  
  content: String!  
  type: MemoryType!  
  source: String!  
  created: DateTime!  
  modified: DateTime!  
  creator: DID!  
  metadata: JSON  
  relations: Relations  
  permissions: Permissions  
}
```

```
type Query {  
  memory(id: ID!): Memory  
  memories(  
    filter: MemoryFilter  
    sort: MemorySort  
    limit: Int = 20  
    offset: Int = 0  
  ): MemoryConnection!  
  
  search(  
    query: String!  
    filters: SearchFilters  
    limit: Int = 20  
  ): [SearchResult!]!  
}
```

```
type Mutation {  
  createMemory(input: CreateMemoryInput!): Memory!  
  updateMemory(id: ID!, input: UpdateMemoryInput!): Memory!  
  deleteMemory(id: ID!): Boolean!  
  
  shareMemory(id: ID!, with: DID!, permissions: [Permission!]!): Memory!  
  revokeAccess(id: ID!, from: DID!): Memory!  
}
```

```
type Subscription {  
  memoryCreated(userId: DID!): Memory!  
  memoryUpdated(userId: DID!): Memory!
```

```
memoryShared(userId: DID!): SharedMemoryEvent!  
}
```

Security & Privacy

Encryption Architecture

At-Rest Encryption

```
javascript
```

```
class MemoryEncryption {
  async encryptMemory(memory, userKey) {
    // Generate memory-specific key
    const memoryKey = await this.deriveKey(userKey, memory.id);

    // Encrypt content
    const encryptedContent = await this.encrypt(
      memory.core.content,
      memoryKey
    );

    // Encrypt metadata selectively
    const encryptedMetadata = await this.encryptMetadata(
      memory.metadata,
      memoryKey
    );

    return {
      ...memory,
      core: {
        ...memory.core,
        content: encryptedContent,
        encrypted: true
      },
      metadata: encryptedMetadata
    };
  }

  async encrypt(data, key) {
    const iv = crypto.getRandomValues(new Uint8Array(12));
    const encrypted = await crypto.subtle.encrypt(
      {
        name: 'AES-GCM',
        iv: iv
      },
      key,
      new TextEncoder().encode(data)
    );

    return {
      ciphertext: btoa(String.fromCharCode(...new Uint8Array(encrypted))),
      iv: btoa(String.fromCharCode(...iv))
    };
  }
}
```

```
}  
}
```

Zero-Knowledge Architecture

javascript

```
class ZeroKnowledgeStore {  
  async store(memory, userPassword) {  
    // Client-side encryption  
    const key = await this.deriveKeyFromPassword(userPassword);  
    const encrypted = await this.encrypt(memory, key);  
  
    // Server never sees plaintext  
    const response = await fetch('/api/store', {  
      method: 'POST',  
      body: JSON.stringify({  
        encrypted: encrypted,  
        proof: this.generateProof(encrypted)  
      })  
    });  
  
    return response.json();  
  }  
  
  generateProof(encrypted) {  
    // Zero-knowledge proof that data is valid  
    // without revealing content  
    return zkSnark.prove({  
      public: [encrypted.hash],  
      private: [encrypted.content],  
      circuit: this.validationCircuit  
    });  
  }  
}
```

Privacy Controls

Data Minimization

javascript

```
class PrivacyFilter {
  sanitizeMemory(memory, level = 'standard') {
    const filters = {
      minimal: ['content', 'type', 'created'],
      standard: ['content', 'type', 'created', 'source', 'basic_metadata'],
      full: null // No filtering
    };

    if (level === 'full') return memory;

    const allowedFields = filters[level];
    return this.pickFields(memory, allowedFields);
  }

  anonymizeMemory(memory) {
    return {
      ...memory,
      header: {
        ...memory.header,
        creator: this.hashDID(memory.header.creator)
      },
      metadata: this.stripPII(memory.metadata)
    };
  }
}
```

Federation Network

Peer-to-Peer Protocol

Node Discovery

```
javascript
```



```
class NodeDiscovery {
  async discoverPeers() {
    const methods = [
      this.dnsDiscovery(),    // DNS TXT records
      this.dhtDiscovery(),    // Distributed Hash Table
      this.mdnsDiscovery(),   // Local network
      this.bootstrapNodes()   // Known bootstrap nodes
    ];

    const peers = await Promise.all(methods);
    return this.deduplicate(peers.flat());
  }

  async dnsDiscovery() {
    const txtRecords = await dns.resolveTxt('_hml._tcp.example.com');
    return txtRecords.map(record => this.parsePeerInfo(record));
  }

  async dhtDiscovery() {
    const dht = new DHT({ bootstrap: this.bootstrapNodes });
    return await dht.findPeers('hml:network');
  }
}
```

Replication Protocol

javascript

```

class ReplicationProtocol {
  async replicate(memory, targetNodes = 3) {
    // Select nodes based on criteria
    const nodes = await this.selectNodes({
      geography: 'distributed',
      reputation: 'high',
      capacity: 'available',
      count: targetNodes
    });

    // Shard memory for redundancy
    const shards = this.createShards(memory, nodes.length);

    // Distribute shards
    const results = await Promise.all(
      shards.map((shard, i) =>
        this.sendShard(shard, nodes[i])
      )
    );

    // Store replication map
    await this.storeReplicationMap(memory.id, results);

    return results;
  }

  createShards(memory, n) {
    // Reed-Solomon erasure coding
    const encoder = new ReedSolomon(n, Math.floor(n * 0.6));
    return encoder.encode(memory);
  }
}

```

Identity & Trust

Decentralized Identity (DID)

javascript

```
class DIDManager {
  async createDID(method = 'web') {
    const methods = {
      web: this.createWebDID,
      key: this.createKeyDID,
      ethr: this.createEthereumDID,
      ion: this.createIONDID
    };

    const did = await methods[method]();
    const document = await this.createDIDDocument(did);

    return {
      did: did,
      document: document,
      keys: await this.generateKeys(did)
    };
  }

  async createDIDDocument(did) {
    return {
      '@context': 'https://www.w3.org/ns/did/v1',
      id: did,
      authentication: [{
        id: `${did}#keys-1`,
        type: 'Ed25519VerificationKey2020',
        controller: did,
        publicKeyMultibase: await this.getPublicKey()
      }],
      service: [{
        type: 'HMLService',
        serviceEndpoint: 'https://hml.example.com/api'
      }]
    };
  }
}
```

Development Guide

Getting Started

Installation

```
bash
```

```
# Install Emma CLI
```

```
npm install -g @emma-hml/cli
```

```
# Initialize new project
```

```
emma init my-memory-app
```

```
# Install SDK
```

```
npm install @emma-hml/sdk
```

Basic Implementation

```
javascript
```

// JavaScript/TypeScript SDK

```
import { HML, MTAP, MCP } from '@emma-hml/sdk';
```

// Initialize HML

```
const hml = new HML({  
  storage: 'local', // local|edge|federation  
  encryption: true,  
  did: 'did:web:user.example.com'  
});
```

// Create memory

```
const memory = await hml.createMemory({  
  content: 'Important information to remember',  
  type: 'text',  
  metadata: {  
    tags: ['important'],  
    source: 'manual'  
  }  
});
```

// Search memories

```
const results = await hml.search('important information', {  
  limit: 10,  
  timeRange: 'last-week'  
});
```

// Get context for AI

```
const context = await hml.getContext({  
  query: 'What did I learn about quantum computing?',  
  maxTokens: 2000  
});
```

Python SDK

python

```
from emma_hml import HML, MemoryStore
import asyncio

class MyMemoryApp:
    def __init__(self):
        self.hml = HML(
            storage_type='local',
            user_id='did:web:user.example.com'
        )

    async def save_conversation(self, messages):
        """Save a conversation to memory layer"""
        for message in messages:
            await self.hml.create_memory(
                content=message['content'],
                role=message['role'],
                metadata={
                    'conversation_id': message['conversation_id'],
                    'timestamp': message['timestamp']
                }
            )

    async def get_relevant_context(self, query):
        """Retrieve relevant memories for AI context"""
        memories = await self.hml.search(
            query=query,
            limit=20,
            relevance_threshold=0.7
        )

        return self.format_for_ai(memories)
```

AI Agent Integration

LangChain Integration

```
python
```

```

from langchain.memory import BaseMemory
from emma_hml import HML

class HMLMemory(BaseMemory):
    """LangChain memory implementation using HML"""

    def __init__(self, hml_config):
        self.hml = HML(**hml_config)
        self.memory_key = "hml_context"

    @property
    def memory_variables(self):
        return [self.memory_key]

    def load_memory_variables(self, inputs):
        query = inputs.get("input", "")
        memories = self.hml.search(
            query=query,
            limit=10
        )

        context = self._format_memories(memories)
        return {self.memory_key: context}

    def save_context(self, inputs, outputs):
        self.hml.create_memory(
            content=f"User: {inputs['input']}\nAI: {outputs['output']}",
            type="conversation",
            metadata={
                "timestamp": datetime.now().isoformat(),
                "model": "gpt-4"
            }
        )

```

Deployment Strategies

Deployment Architectures

1. Standalone (Local-First)

```

yaml

```

```
# docker-compose.yml
version: '3.8'
services:
  emma-local:
    image: emma-hml/local:latest
    volumes:
      - ./data:/app/data
      - ./config:/app/config
    ports:
      - "8080:8080"
    environment:
      - STORAGE_MODE=local
      - ENCRYPTION=true
      - FEDERATION=disabled
```

2. Edge Deployment (Cloudflare Workers)

```
javascript
```



```
// wrangler.toml
name = "emma-edge"
type = "javascript"
account_id = "your-account-id"
workers_dev = true

[durable_objects]
bindings = [
  { name = "MEMORY_STORE", class_name = "MemoryStore" }
]

[[kv_namespaces]]
binding = "MEMORY_INDEX"
id = "your-kv-namespace-id"

// worker.js
export default {
  async fetch(request, env) {
    const url = new URL(request.url);

    if (url.pathname.startsWith('/api/memory')) {
      const id = env.MEMORY_STORE.idFromName(userId);
      const stub = env.MEMORY_STORE.get(id);
      return stub.fetch(request);
    }

    return new Response('Not Found', { status: 404 });
  }
}
```

3. Federation Node

```
yaml
```

```
# kubernetes deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hml-federation-node
spec:
  replicas: 3
  selector:
    matchLabels:
      app: hml-node
  template:
    metadata:
      labels:
        app: hml-node
    spec:
      containers:
        - name: hml-node
          image: emma-hml/federation:latest
          ports:
            - containerPort: 8080 # API
            - containerPort: 9090 # P2P
            - containerPort: 4001 # IPFS
          env:
            - name: NODE_TYPE
              value: "federation"
            - name: ENABLE_IPFS
              value: "true"
            - name: ENABLE_DHT
              value: "true"
```

Performance Optimization

Caching Strategy

```
javascript
```

```
class MemoryCache {
  constructor() {
    this.l1 = new Map();    // In-memory cache
    this.l2 = new Redis();  // Redis cache
    this.l3 = new CDN();    // CDN cache
  }

  async get(key) {
    // L1 Cache (Memory)
    if (this.l1.has(key)) {
      return this.l1.get(key);
    }

    // L2 Cache (Redis)
    const l2Result = await this.l2.get(key);
    if (l2Result) {
      this.l1.set(key, l2Result);
      return l2Result;
    }

    // L3 Cache (CDN)
    const l3Result = await this.l3.get(key);
    if (l3Result) {
      await this.l2.set(key, l3Result);
      this.l1.set(key, l3Result);
      return l3Result;
    }

    // Origin
    const result = await this.fetchFromOrigin(key);
    await this.setCaches(key, result);
    return result;
  }
}
```

Appendices

A. Protocol Version History

Version	Date	Changes
0.1.0	2024-01	Initial protocol draft
0.5.0	2024-06	Added federation support
0.8.0	2024-10	MCP integration
1.0.0	2025-01	Production release

B. Reference Implementations

- 1. **Emma Lite** - Chrome Extension (14-day MVP)
 - Repository: github.com/emma-hml/emma-lite
 - Language: JavaScript
 - Storage: IndexedDB
- 2. **Emma Core** - Full Implementation
 - Repository: github.com/emma-hml/emma-core
 - Language: TypeScript/Rust
 - Storage: Multi-tier
- 3. **HML SDK** - Developer Kit
 - Repository: github.com/emma-hml/sdk
 - Languages: JS, Python, Go, Rust
 - Documentation: docs.emma-hml.com

C. Standards Compliance

- **W3C DID**: Decentralized Identifiers v1.0
- **IPFS**: InterPlanetary File System
- **OAuth 2.0**: Authorization Framework
- **JWT**: JSON Web Tokens (RFC 7519)
- **OpenAPI**: 3.0 Specification

D. Performance Benchmarks

Operation	Target	Actual	Notes
Memory Creation	< 100ms	45ms	Local storage
Search (1000 memories)	< 200ms	120ms	Text search
Vector Search (10k)	< 500ms	380ms	With embeddings
Federation Sync	< 2s	1.5s	3 nodes
Encryption/Decryption	< 50ms	35ms	AES-256

E. Glossary

- **DID:** Decentralized Identifier
- **DHT:** Distributed Hash Table
- **HML:** Human Memory Layer
- **MCP:** Model Context Protocol
- **MTAP:** Memory Transfer & Access Protocol
- **PII:** Personally Identifiable Information
- **P2P:** Peer-to-Peer
- **ZKP:** Zero-Knowledge Proof

F. Contact & Support

- **Website:** <https://emma-hml.com>
- **Documentation:** <https://docs.emma-hml.com>
- **GitHub:** <https://github.com/emma-hml>
- **Discord:** <https://discord.gg/emma-hml>
- **Email:** support@emma-hml.com

License

This specification is released under the MIT License.

MIT License

Copyright (c) 2025 Emma HML Team

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.